

CS 146 Discussion Week 1

Linear Algebra

Outline

- Linear regression
- Convex functions

Linear Regression

- Input: $\mathbf{x}^{(i)} \in \mathbb{R}^d$ input features/attributes for the i -th example
- Output: $y^{(i)} \in \mathbb{R}$ corresponding label for the i -th example
- Goal: find parameters $\boldsymbol{\theta} \in \mathbb{R}^{d+1}$

such that least squares loss function is minimized.

$$J(\boldsymbol{\theta}) = \frac{1}{2n} \sum_{i=1}^n (h_{\boldsymbol{\theta}}(\mathbf{x}^{(i)}) - y^{(i)})^2$$

$$\text{where } h_{\boldsymbol{\theta}}(\mathbf{x}) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \cdots + \theta_d x = \boldsymbol{\theta}^T \mathbf{x}$$

- How do we minimize it? Gradient descent.

Derivation

- Derivation of $J(\boldsymbol{\theta})$

$$\begin{aligned}\frac{\partial}{\partial \theta_j} J(\boldsymbol{\theta}) &= \frac{\partial}{\partial \theta_j} \frac{1}{2n} \sum_{i=1}^n \left(h_{\boldsymbol{\theta}}(\mathbf{x}^{(i)}) - y^{(i)} \right)^2 \\ &= \frac{\partial}{\partial \theta_j} \frac{1}{2n} \sum_{i=1}^n \left(\sum_{k=0}^d \theta_k x_k^{(i)} - y^{(i)} \right)^2 \\ &= \frac{1}{n} \sum_{i=1}^n \left(\sum_{k=0}^d \theta_k x_k^{(i)} - y^{(i)} \right) \times \frac{\partial}{\partial \theta_j} \left(\sum_{k=0}^d \theta_k x_k^{(i)} - y^{(i)} \right) \\ &= \frac{1}{n} \sum_{i=1}^n \left(\sum_{k=0}^d \theta_k x_k^{(i)} - y^{(i)} \right) x_j^{(i)}\end{aligned}$$

- Vectorization

$$\frac{\partial}{\partial \boldsymbol{\theta}} J(\boldsymbol{\theta}) = \frac{1}{n} \sum_{i=1}^n (\boldsymbol{\theta}^T \mathbf{x}^{(i)} - y^{(i)}) \mathbf{x}^{(i)}$$

Comparing analytical and numerical solutions

```
x = np.array([[1, 2]]).transpose()
y = np.array([[3]])
def J(theta):
    return (theta.transpose() @ x - y) ** 2
```

```
theta = np.array([[3, 4]]).transpose()
print(theta)
```

```
[[3]
 [4]]
```

```
print(J(theta))
```

```
[[64]]
```

```
delta = np.array([[0.000001, 0]]).transpose()
print(((J(theta + delta) - J(theta)) / delta[0]))
delta = np.array([[0, 0.000001]]).transpose()
print(((J(theta + delta) - J(theta)) / delta[1]))
```

```
[[16.00000101]]
[[32.000004]]
```

```
def nabla_J(theta):
    return (theta.transpose() @ x - y) * x
```

```
print(nabla_J(theta))
```

```
[[ 8]
 [16]]
```


Notebook: <https://colab.research.google.com/drive/1LjtgnfjJOYXDSjt94HezebKAXtqOIGo4?usp=sharing>

Exercise

- For function $J(\theta) = \frac{1}{2n} \sum_{i=1}^n (h_{\theta}(\mathbf{x}^{(i)}) - y^{(i)})^2$, compute first-order derivative for it and rewrite the results in vectorization formulation.

Convex Functions

- **Definition:** A real-valued function f is **convex** if the line segment joining any two points lies above the function



- A function f is convex iff for all $0 \leq t \leq 1$ and all $x_1, x_2 \in X$:
$$f(tx_1 + (1-t)x_2) \leq tf(x_1) + (1-t)f(x_2)$$


Visualization

```
# suppose we have 3 data points
# x^{(1)} = [1, 1]^T      y^{(1)} = 1
# x^{(2)} = [1, -1.5]^T   y^{(2)} = 2
# x^{(3)} = [1, 0.5]^T    y^{(3)} = 4
def J(theta0, theta1):
    return (theta0 + theta1 - 1) ** 2 + \
           (theta0 - 1.5 * theta1 - 2) ** 2 + \
           (theta0 + 0.5 * theta1 - 4) ** 2
```

```
theta0 = np.linspace(-60, 60, 30)
theta1 = np.linspace(-60, 60, 30)
```

```
theta0, theta1 = np.meshgrid(theta0, theta1)
Js = J(theta0, theta1)
```

```
fig = plt.figure()
ax = plt.axes(projection='3d')
ax.contour3D(theta0, theta1, Z, 50, cmap='binary')
ax.set_xlabel('theta0')
ax.set_ylabel('theta1')
ax.set_zlabel('z');
```



Convex Functions

- Prove that $J(\boldsymbol{\theta}) = \frac{1}{2n} \sum_{i=1}^n (h_{\boldsymbol{\theta}}(\mathbf{x}^{(i)}) - y^{(i)})^2$ is convex.
 - By showing that the second order derivative is positive semi-definite.