

# CS 146 Discussion Week 2

Linear Basis Function Models, Generalization

# Outline

- Linear Basis Function Models
- Generalization
  - Regularization
  - Model Selection
- We will use this notebook during the discussion:
  - <https://colab.research.google.com/drive/1Lur43Fz9lkK8oNvf97GYSQaEsaKPI1rK?usp=sharing>
  - Notebook is adapted from <https://jakevdp.github.io/PythonDataScienceHandbook/05.06-linear-regression.html>

# Linear Basis Function Models

- Generalized class of linear hypothesis

$$h_{\theta}(\mathbf{x}) = \boldsymbol{\theta}^T \boldsymbol{\phi}(\mathbf{x}) = \sum_{j=0}^k \theta_j \phi_j(\mathbf{x})$$

- $\boldsymbol{\phi}(\mathbf{x}): \mathbb{R}^d \rightarrow \mathbb{R}^k$  is a  $k$ -dimensional basis with parameters  $\boldsymbol{\theta} \in \mathbb{R}^k$

- Polynomial basis functions (polynomial regression) with degree  $k$

$$\phi_j(x) = x^j$$

$$h_{\theta}(x) = \sum_{j=0}^k \theta_j \phi_j(x) = \theta_0 + \theta_1 x + \dots + \theta_k x^k$$

- Gaussian basis functions (we use  $k$  of them)

$$\phi_j(x) = e^{-\frac{1}{2}\left(\frac{x-\mu_j}{\sigma}\right)^2}, \phi_0(x) = 1$$

$$h_{\theta}(x) = \sum_{j=0}^k \theta_j \phi_j(x)$$

# Regularized Linear Regression

- Loss Function:  $J(\theta) = \frac{1}{2n} \sum_{i=1}^n \left( h_{\theta}(\mathbf{x}^{(i)}) - y^{(i)} \right)^2 + \frac{\lambda}{2} \sum_{j=1}^d \theta_j^2$
- Gradient Descent Update (note no regularization on bias term)

$$\begin{aligned}\theta_0 &\leftarrow \theta_0 - \alpha \frac{1}{n} \sum_{i=1}^n \left( h_{\theta}(\mathbf{x}^{(i)}) - y^{(i)} \right) \\ \theta_j &\leftarrow \theta_j - \alpha \frac{1}{n} \sum_{i=1}^n \left( h_{\theta}(\mathbf{x}^{(i)}) - y^{(i)} \right) x_j^{(i)} - \alpha \lambda \theta_j\end{aligned}$$

- Closed Form Solution (exercise if you have time: derive out that the inverse always exists for positive  $\lambda$ )

$$\boldsymbol{\theta} = \left( \mathbf{X}^{\top} \mathbf{X} + \lambda \begin{bmatrix} 0 & 0 & 0 & \dots & 0 \\ 0 & 1 & 0 & \dots & 0 \\ 0 & 0 & 1 & \dots & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & 0 & \dots & 1 \end{bmatrix} \right)^{-1} \mathbf{X}^{\top} \mathbf{y}$$

# Linear Regression Exercise

- Given training dataset below, compute the closed-form linear regression solution.

$i$	1	2	3	4	5
$x^{(i)}$	-2	-1	0	1	2
$y^{(i)}$	1	3	5	7	9

- Formula for closed form solution:

$$\boldsymbol{\theta} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

# Linear Regression Exercise

- Put all data points into vectorized form and prepend 1 (for the bias term)

$$\mathbf{X} = \begin{bmatrix} 1 & -2 \\ 1 & -1 \\ 1 & 0 \\ 1 & 1 \\ 1 & 2 \end{bmatrix} \qquad \mathbf{y} = \begin{bmatrix} 1 \\ 3 \\ 5 \\ 7 \\ 9 \end{bmatrix}$$

- Apply formula  $\boldsymbol{\theta} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$
- $\boldsymbol{\theta} = \begin{bmatrix} 5 \\ 2 \end{bmatrix}$

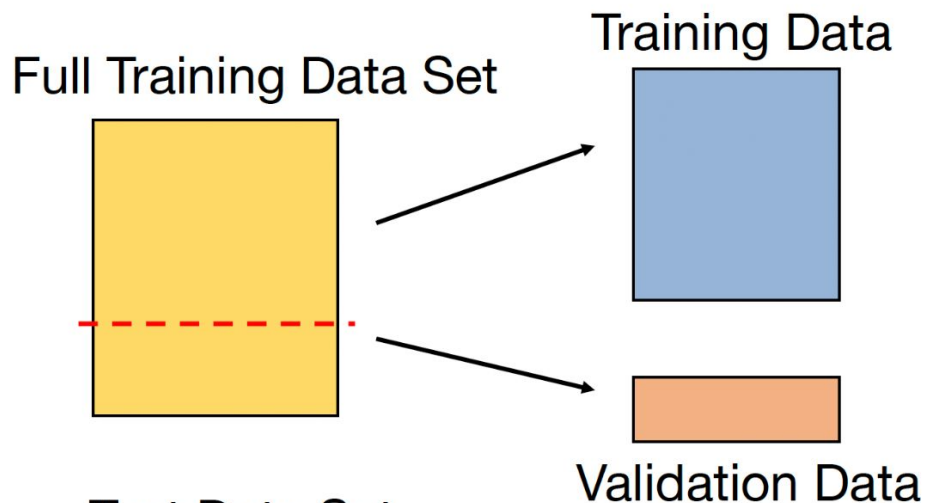
# Regularized Linear Regression Exercise

- Same dataset, compute closed-form solution for regularized linear regression (ridge regression) with  $\lambda = 1$
- Apply formula

$$\boldsymbol{\theta} = \left( \mathbf{X}^T \mathbf{X} + \lambda \begin{bmatrix} 0 & 0 & 0 & \dots & 0 \\ 0 & 1 & 0 & \dots & 0 \\ 0 & 0 & 1 & \dots & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & 0 & \dots & 1 \end{bmatrix} \right)^{-1} \mathbf{X}^T \mathbf{y}$$

- $\boldsymbol{\theta} = \begin{bmatrix} 5 \\ 1.82 \end{bmatrix}$
- The bias term (5) stays the same
- Exercise (if you have time): show that bias term does not change from linear regression to ridge regression. This is true because we do not regularize the bias term.

# Training, Validation, Test data



## Idea:

Train each model on a subset of full training data...

...and then evaluate each model's accuracy on the **validation data** a.k.a. held-out data

Test Data Set



**Intuition:** If validation data  $\sim$  test data, then we have a decent guess for generalization



# Cross-Validation

