

Introduction  
CS 111  
Winter 2023  
Operating System Principles  
Peter Reiher

# Outline

- Administrative materials
- Introduction to the course
  - Why study operating systems?
  - Basics of operating systems

# Administrative Issues

- Instructor and TAs
- Load and prerequisites
- Web site, syllabus, reading, and lectures
- Exams, homework, projects
- Grading
- Academic honesty

# Instructor: Peter Reiher

- UCLA Computer Science department faculty member
- Long history of research in operating systems
- Email: [reiher@cs.ucla.edu](mailto:reiher@cs.ucla.edu)
- Office: No in person office hours this quarter
  - Office hours: TTh 10-11 AM, via Zoom
    - Zoom ID for office hours to be announced later
  - Often available at other times

# TAs

- Section 1A: Victor Zhang
  - Email: [victorzhangyf@ucla.edu](mailto:victorzhangyf@ucla.edu)
- Section 1B: Rustem Can Aygun
  - Email: [canaygun10@gmail.com](mailto:canaygun10@gmail.com)
- Section 1C: Yadi Cao
  - Email: [yadicao95@ucla.edu](mailto:yadicao95@ucla.edu)
- Section 1D: Salekh Parkhati
  - Email: [salekh@cs.ucla.edu](mailto:salekh@cs.ucla.edu)
- Section 1E: Shruthi Srinarasi
  - Email: [shruthi223@g.ucla.edu](mailto:shruthi223@g.ucla.edu)
- Office hours to be announced later

# Instructor/TA Division of Responsibilities

- Instructor handles all lectures, readings, and tests
  - Ask me about issues related to these
- TAs handles projects
  - Ask them about issues related to these
- Generally, instructor won't be involved with project issues
  - So direct those questions to the TAs

# Web Site

- We'll primarily use a web site set up for this class
  - <https://bruinlearn.ucla.edu/courses/153928>
  - Schedules for reading, lectures, exams, projects
  - Project materials
    - And uploads of completed projects
  - Copies of lecture slides
  - Also Zoom IDs for lectures and taped lectures
  - Announcements
  - Sample midterm and final problems

# Prerequisite Subject Knowledge

- CS 32 programming
  - Objects, data structures, queues, stacks, tables, trees
- CS 33 systems programming
  - Assembly language, registers, memory
  - Linkage conventions, stack frames, register saving
- CS 35L Software Construction Laboratory
  - Useful software tools for systems programming
- If you haven't taken these classes, expect to have a hard time in 111



# Course Format

- Two weekly reading assignments
  - Mostly from the primary text
  - Some supplementary materials available on web
- Two weekly lectures
- Four (10-25 hour) individual projects
  - Exploring and exploiting OS features
  - Plus one warm-up project
- A midterm and a final exam

# Course Load

- Reputation: THE hardest undergrad CS class
  - Fast pace through much non-trivial material
- Expectations you should have
  - lectures                      4-6 hours/week
  - reading                        3-6 hours/week
  - projects                        3-20 hours/week
  - exam study                    5-15 hours (twice)
- Keeping up (week by week) is critical
  - Catching up is extremely difficult

# Primary Text for Course

- Remzi and Andrea Arpaci-Dusseau: *Operating Systems: Three Easy Pieces*
  - Freely available on line at  
<http://pages.cs.wisc.edu/~remzi/OSTEP/>
- Supplemented with web-based materials

# Course Grading

- Basis for grading:
  - Class evaluation 1%
  - 1 midterm exam 20%
  - Final exam 26%
  - Lab 0 9%
  - Other labs 11% each
- I do look at distribution for final grades
  - But don't use a formal curve
- All scores available on MyUCLA
  - Please check them for accuracy
  - Scores on BruinLearn not authoritative

# Midterm Examination

- When: 6th week (Tuesday, February 14)
  - Replacing that day's class
  - You can take it online during any two hour period that day
- Scope: All material up to the exam date
  - Approximately 60% lecture, 40% text
  - No questions on purely project materials
- Format:
  - On line, multiple choice, open book/notes
- Goals:
  - Test understanding of key concepts
  - Test ability to apply principles to practical problems

# Final Exam

- When: Friday, March 24
  - You can take it online during any 3 hour period that day
- Scope: Entire course
- Format:
  - On line, multiple choice, open book/notes
- Goals:
  - Determining if you have mastered the full range of material presented in the class

# Lab Projects

- Format:
  - 1 warm-up project
  - 4 regular projects
  - Done individually
- Goals:
  - Develop ability to exploit OS features
  - Develop programming/problem solving ability
  - Practice software project skills

# Late Assignments & Make-ups

- Labs
  - Due dates set by TAs
  - *NOTE: They may change from the dates listed on the syllabus*
  - TAs also sets policy on late assignments
  - The TAs will handle all issues related to labs
    - Ask them, not me
    - Don't expect me to overrule their decisions
- Exams
  - Alternate times or make-ups only possible with prior consent of the instructor
  - If you miss a test, too bad



# Academic Honesty

- It is OK to study with friends
  - Discussing problems helps you to understand them
- It is OK to do independent research on a subject
  - There are many excellent treatments out there
- But all work you submit must be your own
  - Do not write your lab answers with a friend
  - Do not copy another student's work
  - Do not turn in solutions from off the web
  - If you do research on a problem, cite your sources
- I decide when two assignments are too similar
  - And I forward them immediately to the Dean
- If you need help, ask the instructor

# Academic Honesty – Projects

- Do your own projects
  - If you need additional help, ask your TA
- You must design and write all your own code
  - Do not ask others how they solved the problem
  - Do not copy solutions from the web, files or listings
  - Cite any research sources you use
- Protect yourself
  - Do not show other people your solutions
  - Be careful with old listings

# Academic Honesty and the Internet

- You might be able to find existing answers to some of the assignments on line
- Remember, if you can find it, so can we
  - And we have, before
- It IS NOT OK to copy the answers from other people's old assignments
  - People who tried that have been caught and referred to the Office of the Dean of Students
- ANYTHING you get off the Internet must be treated as reference material
  - If you use it, quote it and reference it

# Academic Honesty - Tests

- It shouldn't be necessary to say this, but . . .
- The rules for the tests will be stated before the test
- You must take the test yourself
- You must follow general UCLA academic honesty principles

# Introduction to the Course

- Purpose of course and relationships to other courses
- Why study operating systems?
- What is an operating system?

# What Will CS 111 Do?



- Build on concepts from other courses
  - Data structures, programming languages, assembly language programming, computer architectures, ...
- Prepare you for advanced courses
  - Data bases, data mining, and distributed computing
  - Security, fault-tolerance, high availability
  - Network protocols, computer system modelling
- Provide you with foundation concepts
  - Processes, threads, virtual address space, files
  - Capabilities, synchronization, leases, deadlock

# Why Study Operating Systems?


- Why do we have them, in the first place?
- Why are they important?
- What do they do for us?

# Starting From the Bottom

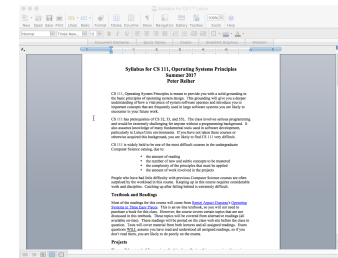
Here's  
what  
you've  
got



Here's  
what  
you  
want

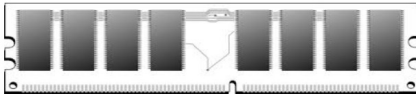


Name	Date Modified	Size	Kind
Lecture_1.pptx	May 11, 2015, 10:11 AM	108 KB	Microsoft PowerPoint Presentation
Lecture_2.pptx	May 11, 2015, 10:11 AM	175 KB	Microsoft PowerPoint Presentation
Lecture_3.pptx	May 11, 2015, 10:11 AM	168 KB	Microsoft PowerPoint Presentation
Lecture_4.pptx	May 11, 2015, 10:11 AM	161 KB	Microsoft PowerPoint Presentation
Lecture_5.pptx	May 11, 2015, 10:11 AM	211 KB	Microsoft PowerPoint Presentation
Lecture_6.pptx	May 11, 2015, 10:11 AM	168 KB	Microsoft PowerPoint Presentation
Lecture_7.pptx	May 11, 2015, 10:11 AM	164 KB	Microsoft PowerPoint Presentation
Lecture_8.pptx	May 11, 2015, 10:11 AM	165 KB	Microsoft PowerPoint Presentation
Lecture_9.pptx	May 11, 2015, 10:11 AM	177 KB	Microsoft PowerPoint Presentation
Lecture_10.pptx	May 11, 2015, 10:11 AM	175 KB	Microsoft PowerPoint Presentation
Lecture_11.pptx	May 11, 2015, 10:11 AM	165 KB	Microsoft PowerPoint Presentation
Lecture_12.pptx	May 11, 2015, 10:11 AM	165 KB	Microsoft PowerPoint Presentation
Lecture_13.pptx	May 11, 2015, 10:11 AM	165 KB	Microsoft PowerPoint Presentation
Lecture_14.pptx	May 11, 2015, 10:11 AM	165 KB	Microsoft PowerPoint Presentation
Lecture_15.pptx	May 11, 2015, 10:11 AM	165 KB	Microsoft PowerPoint Presentation
Lecture_16.pptx	May 11, 2015, 10:11 AM	165 KB	Microsoft PowerPoint Presentation
Lecture_17.pptx	May 11, 2015, 10:11 AM	165 KB	Microsoft PowerPoint Presentation
Lecture_18.pptx	May 11, 2015, 10:11 AM	165 KB	Microsoft PowerPoint Presentation
Lecture_19.pptx	May 11, 2015, 10:11 AM	165 KB	Microsoft PowerPoint Presentation
Lecture_20.pptx	May 11, 2015, 10:11 AM	165 KB	Microsoft PowerPoint Presentation





# What Can You Do With What You've Got?




Read or write some binary words



MOV  
ADD  
JMP  
SQRTPD



Report X and Y  
movements



READ  
REQUEST SENSE



Write to groups  
of pixels



Read or write a block of data

# And You Want This?



# You're Going to Need Some Help

- And that's what the operating system is about
- Helping you perform complex operations
  - That interact
  - Using various hardware
  - And probably various bits of software
- While hiding the complexity
- And making sure nothing gets in the way of anything else

# What Is An Operating System, Anyway?

- System software intended to provide support for higher level applications
  - Including higher level system software applications
  - But primarily for user processes
- The software that sits between the hardware and everything else
- The software that hides nasty details
  - Of hardware, software, and common tasks
- On a good day, the OS is your best computing friend

# But Why Are You Studying Them?

- High probability none of you will ever write an operating system
  - Or even fix an operating system bug
- Not very many different operating systems are in use
  - So the number of developers for them is small
- So why should you care about them?

# Everybody Has One

- Practically every computing device you will ever use has an operating system
  - Servers, laptops, desktop machines, tablets, smart phones, game consoles, set-top boxes
- Many things you don't think of as computers have CPUs inside
  - Usually with an operating system
  - Internet of Things devices
- So you will work with operating systems

# How Do You Work With OSes?

- You configure them
- You use their features when you write programs
- You rely on *services* that they offer
  - Memory management
  - Persistent storage
  - Scheduling and synchronization
  - Interprocess communications
  - Security

# Another Good Reason

- Many hard problems have been tackled in the context of operating systems
  - How to coordinate separate computations
  - How to manage shared resources
  - How to virtualize hardware and software
  - How to organize communications
  - How to protect your computing resources
- The operating system solutions are often applicable to programs and systems you write



# Some OS Wisdom

- View services as objects and operations
  - Behind every object there is a data structure
- Interface vs. implementation
  - An implementation is not a specification
  - Many compliant implementations are possible
  - Inappropriate dependencies cause problems
- An interface specification is a contract
  - Specifies responsibilities of producers & consumers
  - Basis for product/release interoperability

# More OS Wisdom

- Modularity and functional encapsulation
  - Complexity hiding and appropriate abstraction
- Separate policy from mechanism
  - Policy determines what can/should be done
  - Mechanism implements basic operations to do it
  - Mechanisms shouldn't dictate or limit policies
  - Policies must be changeable without changing mechanisms
- Parallelism and asynchrony are powerful and vital
  - But dangerous when used carelessly
- Performance and correctness are often at odds
  - Correctness doesn't always win . . .

# What Is An Operating System?

- Many possible definitions
- One is:
  - It is low level software . . .
  - That provides better, more usable abstractions of the hardware below it
  - To allow easy, safe, fair use and sharing of those resources


# What Does an OS Do?

- It manages hardware for programs
  - Allocates hardware and manages its use
  - Enforces controlled sharing (and privacy)
  - Oversees execution and handles problems
- It abstracts the hardware
  - Makes it easier to use and improves SW portability
  - Optimizes performance
- It provides new abstractions for applications
  - Powerful features beyond the bare hardware

# What Does An OS Look Like?

- A set of management & abstraction services
  - Invisible, they happen behind the scenes
- Applications see objects and their services
  - CPU supports data-types and operations
    - bytes, shorts, longs, floats, pointers, ...
    - add, subtract, copy, compare, indirection, ...
  - So does an operating system, but at a higher level
    - files, processes, threads, devices, ports, ...
    - create, destroy, read, write, signal, ...
- An OS extends a computer
  - Creating a much richer virtual computing platform
    - Supporting richer objects, more powerful operations

# Where Does the OS Fit In?



# What's Special About the OS?

- It is always in control of the hardware
  - Automatically loaded when the machine boots
  - First software to have access to hardware
  - Continues running while apps come & go
- It alone has complete access to hardware
  - Privileged instruction set, all of memory & I/O
- It mediates applications' access to hardware
  - Block, permit, or modify application requests
- It is trusted
  - To store and manage critical data
  - To always act in good faith
- If the OS crashes, it takes everything else with it
  - So it better not crash . . .

# Instruction Set Architectures (ISAs)

- The set of instructions supported by a computer
  - Which bit patterns correspond to what operations
- There are many different ISAs (all incompatible)
  - Different word/bus widths (8, 16, 32, 64 bit)
  - Different features (low power, DSPs, floating point)
  - Different design philosophies (RISC vs. CISC)
  - Competitive reasons (x86, ARM, PowerPC)
- They usually come in families
  - Newer models add features (e.g., Pentium vs. 386)
  - But remain upwards-compatible with older models
    - A program written for an ISA will run on any compliant CPU



# Privileged vs. General Instructions

- Most modern ISAs divide the instruction set into *privileged* vs. *general*
- Any code running on the machine can execute general instructions
- Processor must be put into a special mode to execute privileged instructions
  - Usually only in that mode when the OS is running
  - Privileged instructions do things that are “dangerous”

# Platforms

- ISA doesn't completely define a computer
  - Functionality beyond user mode instructions
    - Interrupt controllers, DMA controllers
    - Memory management unit, I/O busses
    - BIOS, configuration, diagnostic features
    - Multi-processor & interconnect support
  - I/O devices
    - Display, disk, network, serial device controllers
- These variations are called “platforms”
  - The platform on which the OS must run
  - There are lots of them

# Portability to Multiple ISAs

- A successful OS will run on many ISAs
  - Some customers cannot choose their ISA
  - If you don't support it, you can't sell to them
- Which implies that the OS will abstract the ISA
- Minimal assumptions about specific HW
  - General frameworks are HW independent
    - File systems, protocols, processes, etc.
  - HW assumptions isolated to specific modules
    - Context switching, I/O, memory management
  - Careful use of types
    - Word length, sign extension, byte order, alignment
- How can an OS manufacturer distribute to all these different ISAs and platforms?

# Binary Distribution Model

- Binary is the derivative of source
  - The OS is written in source
  - But a source distribution must be compiled
  - A binary distribution is ready to run
- OSes usually distributed in binary
- One (or more) binary distributions per ISA
- Binary model for platform support
  - Device drivers can be added, after-market
    - Can be written and distributed by 3<sup>rd</sup> parties
    - Same driver works with many versions of OS

# Binary Configuration Model

- Good to eliminate manual/static configuration
  - Enable one distribution to serve all users
  - Improve both ease of use and performance
- Automatic hardware discovery
  - Self-identifying busses
    - PCI, USB, PCMCIA, EISA, etc.
  - Automatically find and load required drivers
- Automatic resource allocation
  - Eliminate fixed sized resource pools
  - Dynamically (re)allocate resources on demand

# What Functionality Is In the OS?

- As much as necessary, as little as possible
  - OS code is very expensive to develop and maintain
- Functionality must be in the OS if it ...
  - Requires the use of privileged instructions
  - Requires the manipulation of OS data structures
  - Must maintain security, trust, or resource integrity
- Functions should be in libraries if they ...
  - Are a service commonly needed by applications
  - Do not actually have to be implemented inside OS
- But there is also the performance excuse
  - Some things may be faster if done in the OS

# Conclusion

- Understanding operating systems is critical to understanding how computers work
- Operating systems interact directly with the hardware
- Operating systems rely on stable interfaces