


03 - Data Link Layer

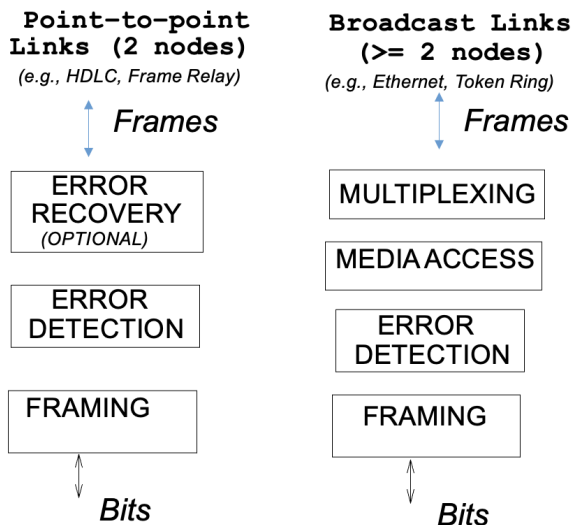
Role of Physical Layer Now

- abstracted to a semi-reliable 1-hop bit-pipe (modem to modem)
- bits - transmitted at physical layer
- frame - transmitted at data link layer (ip packet with ethernet header)
- packet - transmitted at ip/routing layer (tcp packet with ip header)



Data Link Sublayers

- quasi-reliable 1-hop frame-pipe (router to router)
- EOD, output from data link is a frame - a group of bits



- **Framing**: breaking up a stream of bits into units called frames so that we can add extra information like destination addresses and checksums to frames. (Required.)
- **Error detection**: using extra redundant bits called checksums to detect whether any bit in the frame was received incorrectly. (Required).
- **Media Access**: multiple senders. Need traffic control to decide who sends next. (Required for broadcast links).
- **Multiplexing**: Allowing multiple clients to use Data Link. Need some info in the frame header to identify the client. (Optional)
- **Error Recovery**: Go beyond error detection and take recovery action by retransmitting when frames are lost or corrupted. (Optional)

- not usually done in modern routers, assumption is already that not all routers do error recovery, so you can't trust hop-to-hop error recovery
- but modern storage area networks implement hop-to-hop error recovery to ensure low-latency recovery end-to-end by implementing recovery at each hop

Data Framing

Why

- frames allow multiplexing and prevent infinite streams
- frames allow for better error detection and recovery

How

- flag and encoding (HDLC) - add a flag (bit pattern) to delimit frame boundary and encode data to ensure the flag is not preemptively found in the data
- flag and char count (DDCMP) - add flags and a character count, only look for flag after char count
- physical layer flag - supply a special symbol from physical layer to delimit

Fixed-Length Framing


- good for receiving router but bad for variable size frames
- usually used within router code to fragment large payloads

Length-based Framing

- variable length frame with length pre-pended
- still needs a flag to demarcate beginning
- bad if data corrupted b/c requires reading and must be done while reading transmission
- e.g., DECnet, DDCMP

Sentinel-based Framing

- variable length frames with flags delimit at both beginning and end
- add stuffing (stuff some bits) where the flag is found in the data/frame
- receiver "unstuffs" data e.g., sps flag is 111, then any time we see 111 in the data, add a 0 after → 111...1110...111
- irl use byte-stuffing and denote the escape char to prevent false flag assign byte `DLE` and stuff it whenever STX/ETX found in payload if you get `DLE` in data then do `DLE_DLE`



- denote flags as STX/ETX (start/end)

Physical Layer Solution

- because 4-5 encoding produces 16 possible symbols but there are 32 possible encoded, pass 2 of those unused as SOF/EOF

Error Detection

- B/c TCP doesn't require end-to-end checksums, data link undetected errors must be so small close to 1 undetected error per 20 years of data

Types of Errors

- **Random Errors.** A noise spike or inter-symbol interference makes you think a 0 is a 1 or 1 to 0. Fiber: 1 in 10¹⁰
- **Burst errors:** A group of bits get corrupted because of synchronization or connector plugged in. Correlated!
- **Modeling Burst error:** Burst error of length k à distance from first to last is $k - 1$. Intermediate may or may not be corrupted. Burst error of 5 starting at 50. Bits 50 and 54 are corrupted, bits 51-53 may or may not be corrupted
- **Goal for quasi-reliability:** Like to add checksums to detect as large a burst (say 32) and as many random (at least 3)
- **Comparison:** Imagine a frame of size 1000 and an error rate of 1 in 1000. If random, all frames corrupted on average. If we get a burst of 1000 every 1000 frames, only 1 is lost!

Parity error detection to Checksums

- parity - parity of the number of 1s in the data
- doing XOR of bits to detect parity may detect up to 2 bit error, but how to check error for >3 bits
- instead use checksums
- goal is detection not correction, detection = some bit in frame bad so drop frame vs flip bit (correction)
- CRC32 - use mod 2 division (XORs) for checksum instead of sum


Simple Divide Checksum

- Consider message M and generator G to be binary integers.
- Let r be number of bits in G . We find the remainder t of $2^r M$ when divided by G . Why not just M ? So that we can separate checksum from message at receiver by looking at last r bits.
- Thus $2^r M = kG + t$. Thus:
 $2^r M + G - t = (k + 1)G$. So we add a checksum $c = G - t$ to the shifted message and the result should divide G .
Example: $M = 110010$ (50), $r = 3$, $G = 7$. After shifting 3 bits, we get 400. Remainder $t = 1$, checksum = 6.
So we send 110010 110 which is divisible by 7 (406)
- Has reasonable properties. However integer division hard to implement. Prefer to do without carries.

Mod 2 Checksum - CRC (Cyclic Redundancy Check)

- **No carries.** Repeated addition does not result in multiplication. e.g. $1100 + 1100 = 0000$;
 $1100 + 1100 + 1100 = 1100$
- Multiplication is normal except for no carries: e.g. $1001 * 11 = 10010 + 1001 = 11011$. Shift and Ex-or instead of Shift and Add as in normal arithmetic.
- Similar algorithm to ordinary division. Again let r be number of bits in G . We find the remainder c of $2^{r-1}M$ when divided by G . Why **only shift $r - 1$ bits** this time?
- Thus $2^{r-1}M = k.G + c$. Thus $2^{r-1}M - c = k.G$. Thus $2^{r-1}M + c = k.G$ (addition same as subtraction). Send **c as checksum**

- background
- example, observe remainder is less than generator



- For CRC, we need to repeatedly add (mod 2) multiples of the generator until we get a number that is $r - 1$ bits long that is the remainder.
- The only way to reduce number of bits in Mod 2 arithmetic is to remove MSB by adding (mod 2) a number with a 1 in the same position.
- While no more bits
 - If MSB = 1, XOR with generator (**RED**)
 - Shift out MSB and Shift in next bit (**BLUE**)

$$\text{CRC-16: } X^{16} + X^{15} + X^2 + 1 = 11000000000000101$$

We skip proofs of these properties this quarter but they are in your notes, Not required for HWs and tests.

Odd bit errors: can handle but not a big deal as parity can handle with using just 1 bit. 1

Two bit errors specially designed CRCs can do this. Beats parity!

Burst errors: CRC-32 can catch any 32 bit burst error for sure. Further it can catch larger burst errors with very high probability: $(1 - 1/2^{32})$

Summary: So the big deal is that it can for sure catch up to 3 bit errors, and can detect **any** error with very high probability. Like a hash function with with some deterministic guarantees

- CRC is implemented via LFSR in CPU

Error Recovery (Optional)



- usually not done on WAN but done on SAN
- RFC spec - english spec for network protocols

- RFC for error recovery at data link layer must ensure packets delivered without duplication, loss, or mis-ordering

Assumption


- **Assumes error detection:** Assumes undetected error rate small enough to be ignored
- **Loss as well as errors:** whole frames can be lost in a way not detected by error detection
- **FIFO:** Physical layer is FIFO
- **Arbitrary Delay:** Delay on links is arbitrary and can vary from frame to frame.

Time-Space Examples




→ Must defend against early retransmits

- must return ack to validate the sending of the next packet, must id the packets to



→ Need sequence numbers




→ Must ack even duplicates


detect true duplication vs intended duplicate

- [illegible]

Stop and Wait Protocol (Send then wait for ack)



- time state diagram



-----Sender Code-----

Sender keeps state variable **SN**, initially 0 and repeats following loop

- 1) Accept a new packet if available from higher layer and store it in buffer **B**
 - 2). Transmit a frame Send (**SN**, **B**)
 - 3). If error-free (**ACK**, **R**) frame received and $R \neq SN$ then
 $SN = R$
Go to Step 1
- Else if the previous condition does not occur after **T** sec
Go to Step 2

-----Receiver Code-----

Receiver keeps state variable **RN**, initially 0

When an error free data frame (**S**, **D**) is received

On receipt:

If $S = RN$ then

Pass D to higher layer


$RN = RN + 1$;

Send (**ACK**, **RN**) // Send ack unconditionally!

- code for sender and receiver

Band Invariance

- when receiver processes the message in channel and sends ack, all state are only of 1 number (the id of the latest packet acknowledged) → thus we can check for correctness of packets by ensuring band invariance
- prove band invariance by checking state transitions



- 3 other cases: Receive Ack, Send Ack, and Send new frame
- Just need to show that invariant is preserved by these 6 protocol actions/state transitions.

Code of Alternating Bit

-----Sender Code-----

Sender keeps state bit **SN**, initially 0 and repeats following loop

- 1) Accept a new packet if available from higher layer and store it in buffer **B**
- 2). Transmit a frame Send (**SN**, **B**)
- 3). If error-free (**ACK**, **R**) frame received and **R != SN** then
 SN = R
 Go to Step 1
Else if the previous condition does not occur after **T**
 Go to Step 2

-----Receiver Code-----

Receiver keeps state bit **RN**, initially 0

When an error free data frame (**S**, **D**) is received

On receipt:

If **S = RN** then

 Pass D to higher layer

RN = ~ RN ; //flip bit!

 Deliver data m to client.


Send (**ACK**, **RN**) // Send **ack unconditionally!**

- alternating bit recovery code:

Performance Measures

- throughput - jobs completed per second
- latency - worst-case time to complete a job
- 1-way propagation delay - time for the transmitted bit to reach the receiver; disregarding transmission rate, there is some amount of time it takes for the bit to travel the length of the link - this is the 1 way propagation delay
- transmission rate - the rate at which bits can be sent over a link, i.e. the number of bits per second - tells us that the second bit may come quickly after the first bit is sent
- pipe size (bandwidth-delay product) = transmission rate \times round-trip propagation delay

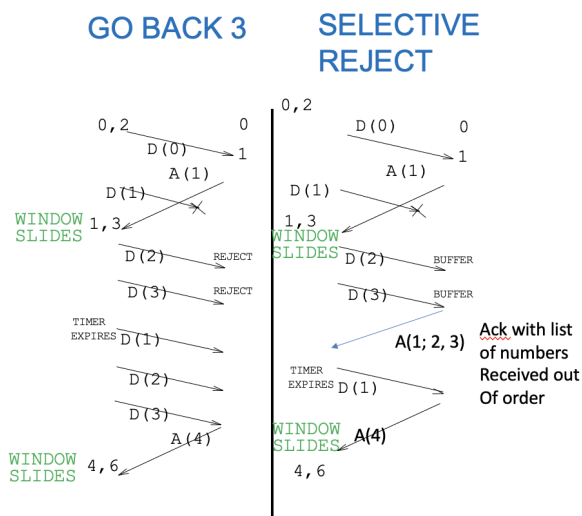
- pipe size (and prop delay) tells us our pipe/link utilization e.g., stop and wait frames
 - 1-way Propagation Delay: 250 msec
 - Transmission speed: 100 kbit/sec
 - Frame size: 1000 bits.
 - What is throughput? 2000 bits per second, which is 2% of a 100,000 bit per second link.



(send next frame after ack)

Sliding Window Protocol

- **Window**: Sender can send a window of outstanding frames before getting any acks. Lower window edge L , can send up to $L + w - 1$.
- **Receiver numbers**: receiver has a receive sequence number R , next number it expects. L and R are initially 0.
- **Sender Code**: Retransmits all frames in current window until it gets an ack. Ack numbered r implicitly acknowledges all numbers $< r$.
- **Two variants**: receiver accepts frames in order only (go-back-N) or buffers out-of-order frames (selective reject)
- we have batched rejects or selective rejects using complex acks or simple acks but resend all packets in the sliding window of frames sent



- code for both implementations

Go Back N Code

Sender Code

Sender keeps state variable L , initially 0

Send (s, m) // send data message m with number s

The sender can send this frame if:

m corresponds to s -th data item

given to sender by client AND

$L \leq s \leq L + w - 1$ // in allowed send window

Receive (r, Ack) // receive an ack number r

On receipt:

$L := r$ // slide lower window edge to ack number

Receiver Code

Receiver keeps state variable R , initially 0

Receive (s, m) // receive data message m with number s

On receipt:

If $s = R$ then

$R := s + 1$

Deliver data m to client.

Send (r, Ack) // send ack with number r

// receivers typically send acks in response to data

// messages but our code can send acks anytime

r must equal R

Selective Reject Sender code

Sender keeps a lower window edge L initially 0 but also an **array** with a bit set for all numbers **acked** so far.

Initially, all bits are clear. In practice, we implement this array by a bitmap of size w which we shift

Send (s, m) // send data message m with number s

The sender can send this frame if:

m corresponds to s -th data item

given to sender by client AND

$L \leq s \leq L + w - 1$ AND

s has not been **acked** // new for selective reject

Receive ($r, \text{List Ack}$) // receive an ack number r with **List**
// of received numbers $> r$

On receipt:

$L := r$ // slide lower window edge to ack number

Mark numbers in **List** as **acked** at sender

Selective Reject Receiver Code

Receiver keeps a receiver number R initially 0 but also an **array** with a bit set for all numbers received so far. Initially, all bits are clear. In practice, we implement this array again by a bitmap of size w which we shift. In addition to the bitmap, we have a buffer for each number where we can store out of order messages

Receive (s, m) // receive data message m with number s

On receipt:

If $s \geq R$ then

Mark s as **acked** and Buffer m 1 0 0 0 0

While R **acked** do

Deliver data message at position R

$R := R + 1$ 1 0 1 0 0

Send ($r, \text{List Ack}$) // send **ack** with number r and **List**
// of received numbers $> r$

r must equal R

List contains received numbers $> R$


- implementation details - ONLY for FIFO packets

Implementation and Other Details

- Timers: works regardless of values, but needed for performance. So calculate round-trip delay.
- Need only one timer (for lowest outstanding number) in Go-back- n . Need one for each window element in Sel Reject.
- In selective reject, have to send an ack with R and a bit-map of numbers greater than R that have been received.
- Piggybacking: to reduce frames sent.

- **Alternating bit:** Modulus is 2 (just one bit)
- **Go back W :** Need a Modulus of $W+1$
- **Selective Reject:** Need a modulus of $2W$

Intuition as to why the window size must be bounded



- flow control - variable size windows, usually set by receiver to prevent backlogging frames (send the right window edge with ack)
- previously selective reject was not allowed as you need long lists/buffers for acks, but now RFC has some allowance

