


12 - Application Protocols

Client-Server vs. P2P

- Client-server - server always on with permanent IP, client connects with dynamic IP
 - e.g., HTTP, IMAP, FTP
- P2P - no such server, all inter client connects, scalable
 - e.g., torrent file sharing - bittorrent

Process Communication

- processes on the same host comms via Inter-Process Comms (IPC) via OS
- client-server - processes are either clients or servers and are always listening or trying to connect
- p2p - procs have both client and server ports



- comms via sockets (abstraction)
- procs identified by both IP and Port
 - port e.g., HTTP:80, SMTP:25 (mail)

Application Level Protocols

- types of messages exchanged,
 - e.g., request, response
 - message syntax:
 - what fields in messages & how fields are delineated
 - message semantics
 - meaning of information in fields
 - rules for when and how processes send & respond to messages
- protocols contain/define
 - open protocols:
 - defined in RFCs, everyone has access to protocol definition
 - allows for interoperability
 - e.g., HTTP, SMTP
 - proprietary protocols:
 - e.g., Skype, Zoom
- apps require transport service
 - data integrity
 - some apps (e.g., file transfer, web transactions) require 100% reliable data transfer
 - other apps (e.g., audio) can tolerate some loss
 - timing
 - some apps (e.g., Internet telephony, interactive games) require low delay to be "effective"
 - throughput
 - some apps (e.g., multimedia) require minimum amount of throughput to be "effective"
 - other apps ("elastic apps") make use of whatever throughput they get
 - security
 - encryption, data integrity, ...

- apps require transport service

application	data loss	throughput	time sensitive?
file transfer/download	no loss	elastic	no
e-mail	no loss	elastic	no
Web documents	no loss	elastic	no
real-time audio/video	loss-tolerant	audio: 5Kbps-1Mbps video: 10Kbps-5Mbps	yes, 10's msec
streaming audio/video	loss-tolerant	same as above	yes, few secs
interactive games	loss-tolerant	Kbps+	yes, 10's msec
text messaging	no loss	elastic	yes and no

- e.g. proc examples

TCP vs. UDP

TCP service:

- **reliable transport** between sending and receiving process
- **flow control**: sender won't overwhelm receiver
- **congestion control**: throttle sender when network overloaded
- **connection-oriented**: setup required between client and server processes
- **does not provide**: timing, minimum throughput guarantee, security

UDP service:

- **unreliable data transfer** between sending and receiving process
 - **does not provide**: reliability, flow control, congestion control, timing, throughput guarantee, security, or connection setup.
- Q: why bother? *Why* is there a UDP?

application	application layer protocol	transport protocol
file transfer/download	FTP [RFC 959]	TCP
e-mail	SMTP [RFC 5321]	TCP
Web documents	HTTP [RFC 7230, 9110]	TCP
Internet telephony	SIP [RFC 3261], RTP [RFC 3550], or proprietary	TCP or UDP
streaming audio/video	HTTP [RFC 7230], DASH	TCP
interactive games	WOW, FPS (proprietary)	UDP or TCP

- e.g., app transport protocols

Vanilla TCP & UDP sockets:

- no encryption
- cleartext passwords sent into socket traverse Internet in cleartext (!)

Transport Layer Security (TLS)

- provides encrypted TCP connections
- data integrity
- end-point authentication

TLS implemented in application layer

- apps use TLS libraries, that use TCP in turn
- cleartext sent into "socket" traverse Internet *encrypted*
- See our earlier lecture or Chapter 8 in Kurose-Ross

- TCP security (TLS)

Web and HTTP

- web page consists of objects stored in a DOM structure
- object can be HTML file/element, images, applets, audio, etc.
- web-page consists of base HTML file (e.g. index.html) with several referenced objects

addressable via URL e.g. www.someschool.edu/someDept/pic.gif

host name
path name

HTTP

- **HTTP** - hypertext transfer protocol

HTTP: hypertext transfer protocol

- Web's application-layer protocol
- client/server model:

- **client**: browser that requests, receives, (using HTTP protocol) and "displays" Web objects
- **server**: Web server sends (using HTTP protocol) objects in response to requests



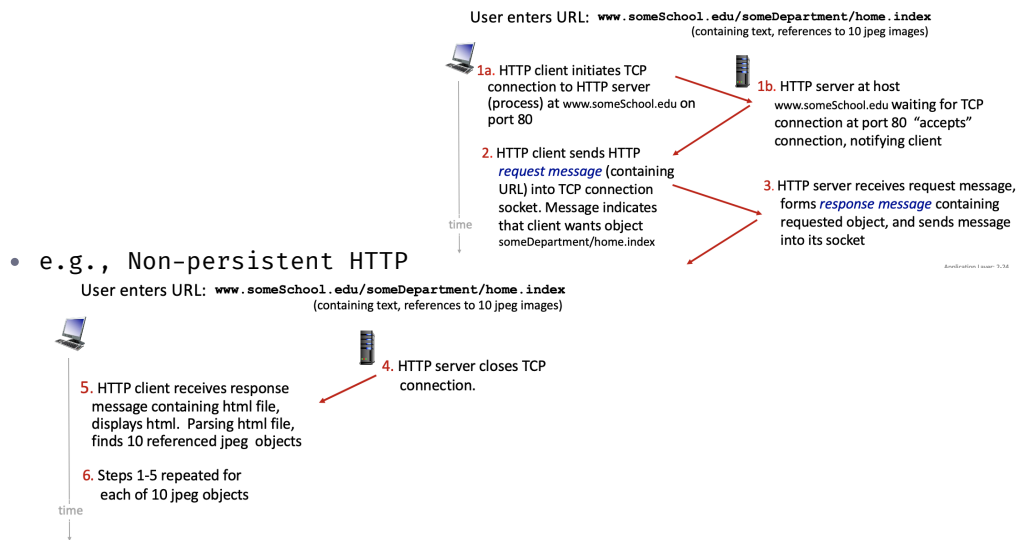
- HTTP is persistent or non-persistent because storing state is complex

Non-persistent HTTP

1. TCP connection opened
 2. at most one object sent over TCP connection
 3. TCP connection closed
- downloading multiple objects required multiple connections

Persistent HTTP

- TCP connection opened to a server
- multiple objects can be sent over *single* TCP connection between client, and that server
- TCP connection closed




- e.g., Non-persistent HTTP

User enters URL: `www.someSchool.edu/someDepartment/home.index`
(containing text, references to 10 jpeg images)

RTT (definition): time for a small packet to travel from client to server and back

HTTP response time (per object):

- one RTT to initiate TCP connection
- one RTT for HTTP request and first few bytes of HTTP response to return
- object/file transmission time



- non-persistent performance

Non-persistent HTTP response time = 2RTT + file transmission time

- persistent (HTTP 1.1) v. non-persistent pro cons


Non-persistent HTTP issues:

- requires 2 RTTs per object
- OS overhead for *each* TCP connection
- browsers often open multiple parallel TCP connections to fetch referenced objects in parallel

Persistent HTTP (HTTP1.1):

- server leaves connection open after sending response
- subsequent HTTP messages between same client/server sent over open connection
- client sends requests as soon as it encounters a referenced object
- as little as one RTT for all the referenced objects (cutting response time in half)

- messages are request (GET/POST/PUT ...) or response (status ode, e.g. 404)



POST method:

- web page often includes form input
- user input sent from client to server in entity body of HTTP POST request message

HEAD method:

- requests headers (only) that would be returned *if* specified URL were requested with an HTTP GET method.

GET method (for sending data to server):

- include user data in URL field of HTTP GET request message (following a '?'):

PUT method:

- uploads new file (object) to server
- completely replaces file that exists at specified URL with content

- request messages

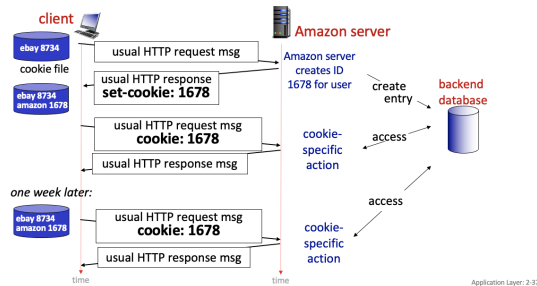
`www.somesite.com/animalsearch?monkeys&banana`

- status code appears in 1st line in server-to-client response message.
- some sample codes:
 - 200 OK
 - request succeeded, requested object later in this message
 - 301 Moved Permanently
 - requested object moved, new location specified later in this message (in Location: field)
 - 400 Bad Request
 - request msg not understood by server
 - 404 Not Found
 - requested document not found on this server
 - 505 HTTP Version Not Supported

- request messages

Cookies - State Management


- HTTP requests are stateless so for multi-step exchange or repeated connections, store cookies to maintain state (stored in the browser)



- components
 - cookie header line of HTTP **response** message
 - cookie header line in next HTTP **request** message
 - cookie kept on user's host and managed by user's browser
 - back-end DB at server website containing cookie id val to map state
- **first party cookies** - track user behavior for the given website
- **third party cookies** - tracks user behavior across multiple websites without visiting third party site
- GDPR (EU general data protection regulation) requires sites to inform users about 3rd party cookies

Web Cache

- to inc performance and decrease load on server, initial http get may include data and web cached data
- may be done via proxy server which acts both as client and server intermediary




- example

Scenario:

- access link rate: 1.54 Mbps
- RTT from institutional router to server: 2 sec
- web object size: 100K bits
- average request rate from browsers to origin servers: 15/sec
 - avg data rate to browsers: 1.50 Mbps

Performance:

- access link utilization = $\frac{1.50}{1.54} \approx .97$ *problem: large queueing delays at high utilization!*
- LAN utilization: .0015
- end-end delay = Internet delay + access link delay + LAN delay = 2 sec + (minutes) + usecs



- og bottleneck


Scenario:

- access link rate: 154 Mbps
- RTT from institutional router to server: 2 sec
- web object size: 100K bits
- average request rate from browsers to origin servers: 15/sec
 - avg data rate to browsers: 1.50 Mbps

Performance:

- access link utilization = $\frac{1.50}{154} \approx .0097$
- LAN utilization: .0015
- end-end delay = Internet delay + access link delay + LAN delay = 2 sec + (msecs) + usecs

Cost: faster access link (expensive!) → msecs




- sol1: get faster link
- sol2: web cache, much cheaper than more expensive link

suppose cache hit rate is 0.4:

- 40% requests served by cache, with low (msec) delay
- 60% requests satisfied at origin
 - rate to browsers over access link = $0.6 * 1.50 \text{ Mbps} = .9 \text{ Mbps}$
 - access link utilization = $0.9 / 1.54 \approx .58$ means low (msec) queueing delay at access link
- average end-end delay:
 - = 0.6 * (delay from origin servers)
 - + 0.4 * (delay when satisfied at cache)
 - = $0.6 (2.01) + 0.4 (\sim \text{msecs}) \approx \sim 1.2 \text{ secs}$

lower average end-end delay than with 154 Mbps link (and cheaper too!)



HTTP/2

- dec delay iin multi obj HTTP requets
- HTTP1.1 introduced multiple pipelined gets over single TCP connection
- server responds in order (FCFS) - this may cause head of line (HOL) blocks for small objs behind large objects → loss recovery stalling transmission

HTTP/2: [RFC 7540, 2015] increased flexibility at server in sending objects to client:


- methods, status codes, most header fields unchanged from HTTP 1.1
- transmission order of requested objects based on client-specified object priority (not necessarily FCFS)
- push unrequested objects to client
- divide objects into frames, schedule frames to mitigate HOL blocking

HTTP 1.1: client requests 1 large object (e.g., video file) and 3 smaller objects



- e.g., HTTP/1.1 HOL issue

HTTP/2: objects divided into frames, frame transmission interleaved



- e.g., HTTP/2 frame sol

HTTP/3


- HTTP/2 disadvantages
 - single connection means recovery from packet loss still stalls object transmissions
 - no security over vanilla TCP
- HTTP/3 adds security and per object error and congestion control via pipelining over UDP

Email, SMTP, IMAP


- composed of: user agents, mail servers, and SMTP
- SMTP** - simple mail transfer protocol
 - between mail servers on client (sender) - server (receiver) protocol

SMTP RFC (5321)

- uses TCP to reliably transfer email message from client (mail server initiating connection) to server, port 25
 - direct transfer: sending server (acting like client) to receiving server
- three phases of transfer
 - SMTP handshaking (greeting)
 - SMTP transfer of messages
 - SMTP closure
- command/response interaction (like HTTP)
 - commands: ASCII text
 - response: status code and phrase




- user agent** - mail reader, creating editing reading sending mail on various clients - outlook, gmail, iphone, etc,
- mail server** - composed of mailbox (contains incoming messages) and message queue (outgoing mail queue)



comparison with HTTP:

- HTTP: client pull
 - SMTP: client push
 - both have ASCII command/response interaction, status codes
 - HTTP: each object encapsulated in its own response message
 - SMTP: multiple objects sent in multipart message
 - SMTP uses persistent connections
 - SMTP requires message (header & body) to be in 7-bit ASCII
 - SMTP server uses CRLF.CRLF to determine end of message
- request and response messages
 - **IMAP** - internet mail access protocol - provides retrieval and deletion



- **SMTP**: delivery/storage of e-mail messages to receiver's server
- mail access protocol: retrieval from server
 - **IMAP**: Internet Mail Access Protocol [RFC 3501]: messages stored on server, IMAP provides retrieval, deletion, folders of stored messages on server
- **HTTP**: gmail, Hotmail, Yahoo!Mail, etc. provides web-based interface on top of SMTP (to send), IMAP (or POP) to retrieve e-mail messages