

# CS/ENGR M148 L14: Backpropagation

Sandra Batista

### **This week in discussion section:**

In lab this week you'll be learning about PyTorch to apply neural networks to your projects.

No project check-ins this week.

Lecture on 11/25/24 will be via Zoom for the holiday.

**Midterm grades posted. Grade distribution on piazza. Regrade requests due by 6 pm 11/20/24.**

2

**PS3 quiz today!**

**Will post final project report guidelines this week also with PS4.**

**Extra credit Final Exam Review Question Code Bank:**

**<https://forms.gle/XdC97wxwWd8QuTR9A>**

**Questions due by 11:59 pm PT on 11/25/24.**

We'll share questions with solutions during week 10 for final exam review.

# Join our slido for the week...

---

<https://app.sli.do/event/nCV57u4mC7eUMit9euSBr2>



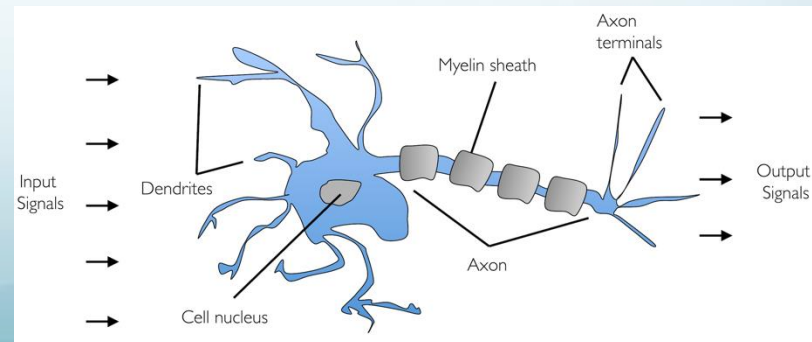
# Today's Learning Objectives

Students will be able to:

- Review: Describe what a neural network (NN) is
- Apply the forward algorithm on the MNIST NN
- Trace Backpropagation on a small NN
- Understand the role of gradient descent in training NN

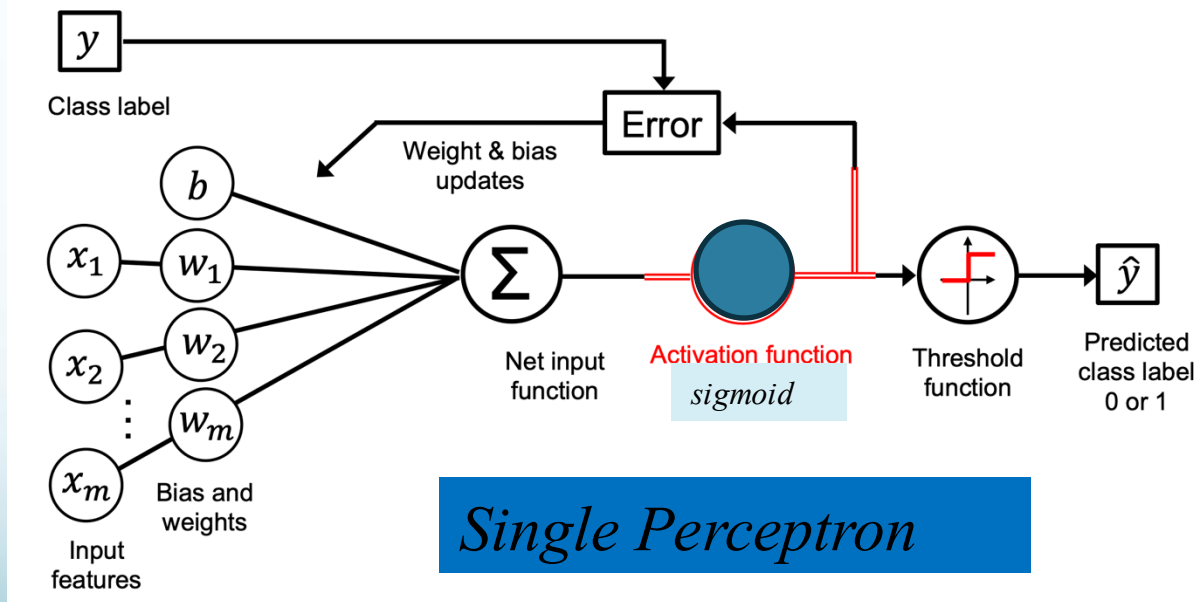
# Neural Networks (NN)

- *Based on models for how brain works using artificial neurons*
- *Many varied successful applications such as mood recognition in pictures, modeling virus mutations, and predicting needed medical resources*
- *Used to model complex, nonlinear models*
- *We'll focus on using them for **classification**.*



# What is a neural network?

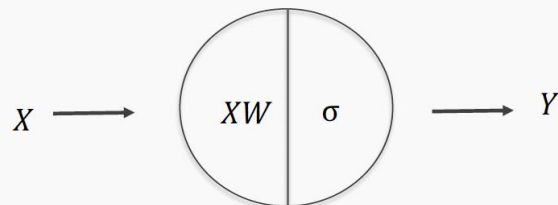
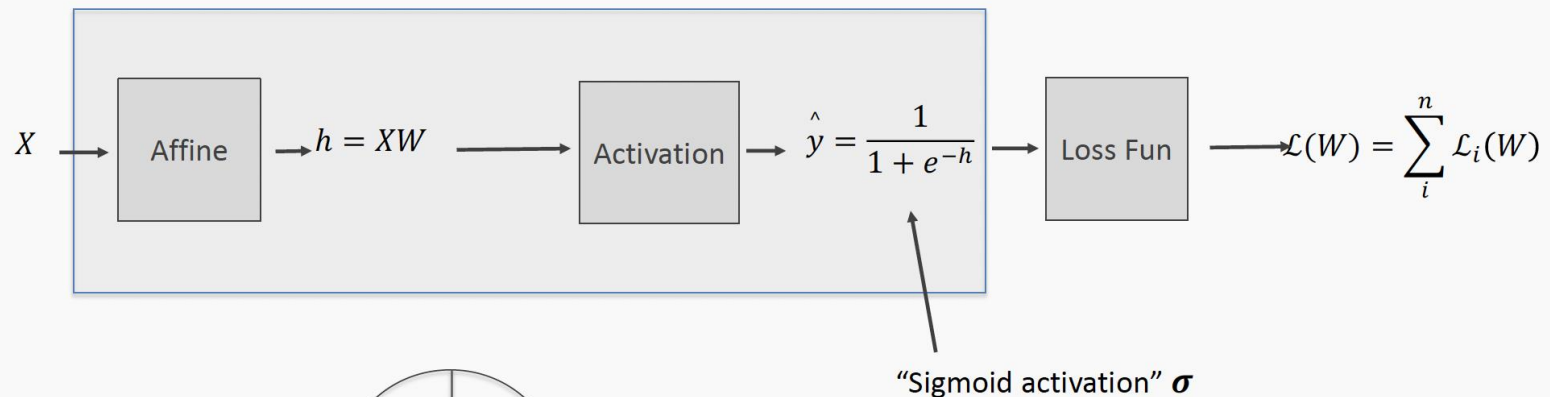
- A neural network consists of layers of nodes or artificial neurons*



A *single* neuron can be a *logistic regression* or linear unit or other activation functions.

# What is a neural network?

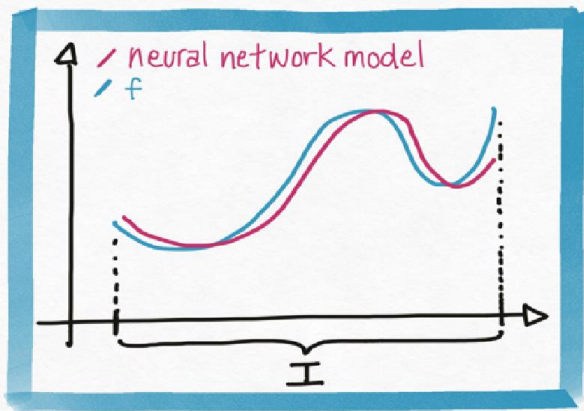
- *A neural network consists of layers of nodes or artificial neurons*



**Single Neuron Network  
aka Perceptron**



# Neural Networks as Universal Approximators



## **Theorem:**

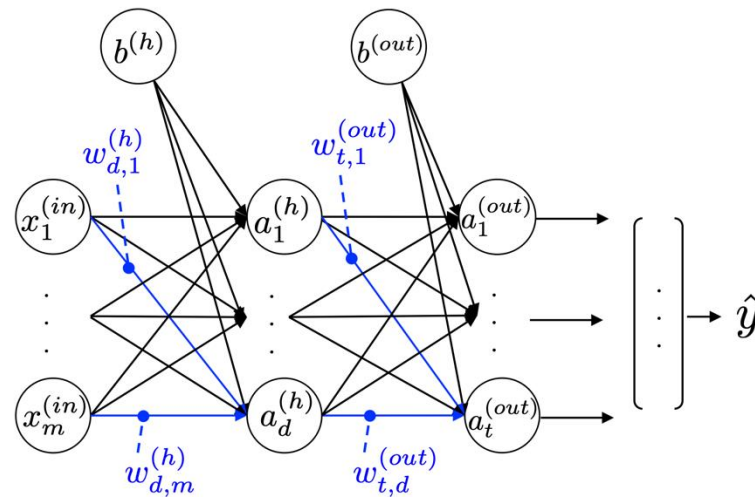
*For any continuous function  $f$  defined on a bounded domain, we can find a neural network that approximates  $f$  with an arbitrary degree of accuracy.*

*One hidden layer is enough to represent an approximation of any function to an arbitrary degree of accuracy.*

*A neural network can **approximate** non-linear functions either for regression or classification.*

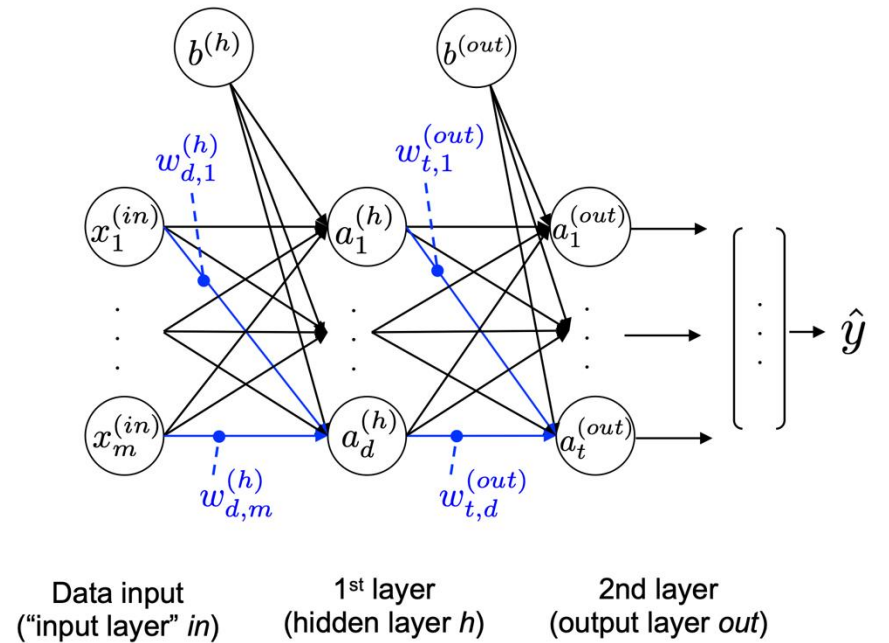
# Multilayer Perceptron

- A neural network is a *combination* of neurons such as logistic regression (or other types) units.
- A **multilayer perceptron (MLP)** is a fully connected network of neurons.
- An MLP is a multilayer **feedforward** NN because each layer is input to the next.
- A network with more than one hidden layer is a **deep NN**.

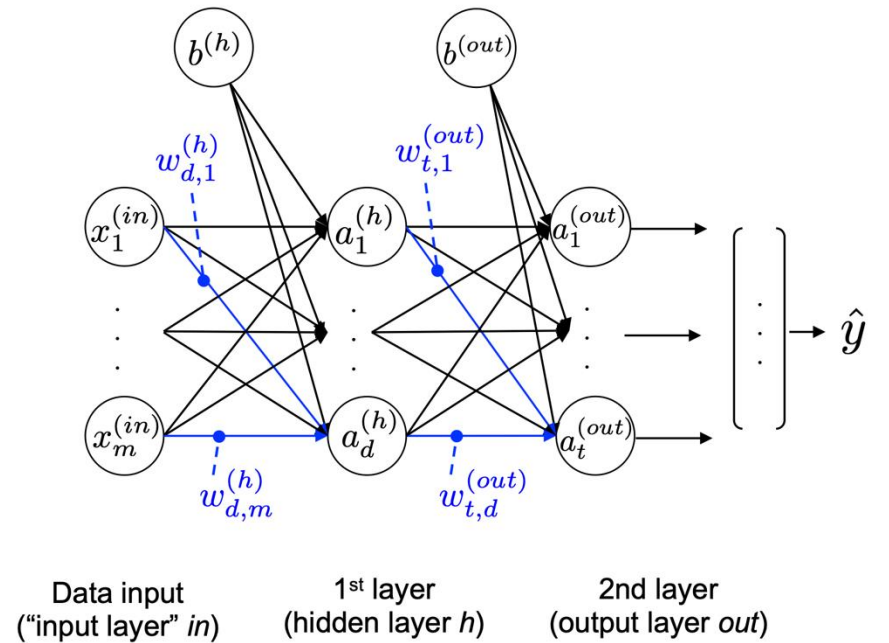


[Raschka et al  
2022]

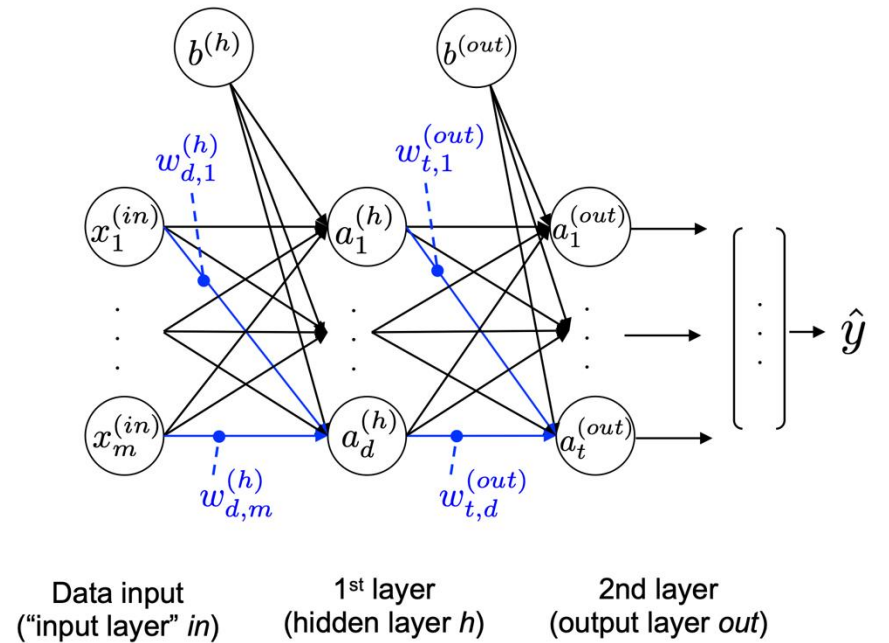
# Multilayer Perceptron



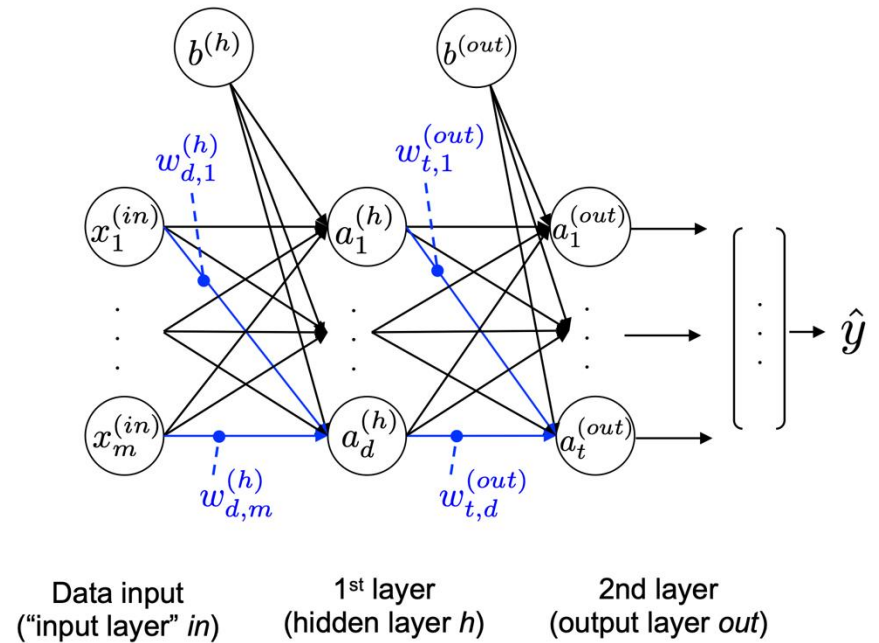
# Multilayer Perceptron



# Multilayer Perceptron



# Composing a complex function



# Activation function

$$h = f(W \cdot X + b)$$

The activation function should:

- Provide **non-linearity**
- Ensure **gradients remain large** through hidden unit

Examples:

- sigmoid
- ReLU
- identity

# Today's Learning Objectives

Students will be able to:

- ✓ Review: Describe what a neural network (NN) is
  - Apply the forward algorithm on the MNIST NN
  - Trace Backpropagation on a small NN
  - Understand the role of gradient descent in training NN



# Modified National Institute of Standards and Technology (MNIST) Classification NN from Scratch

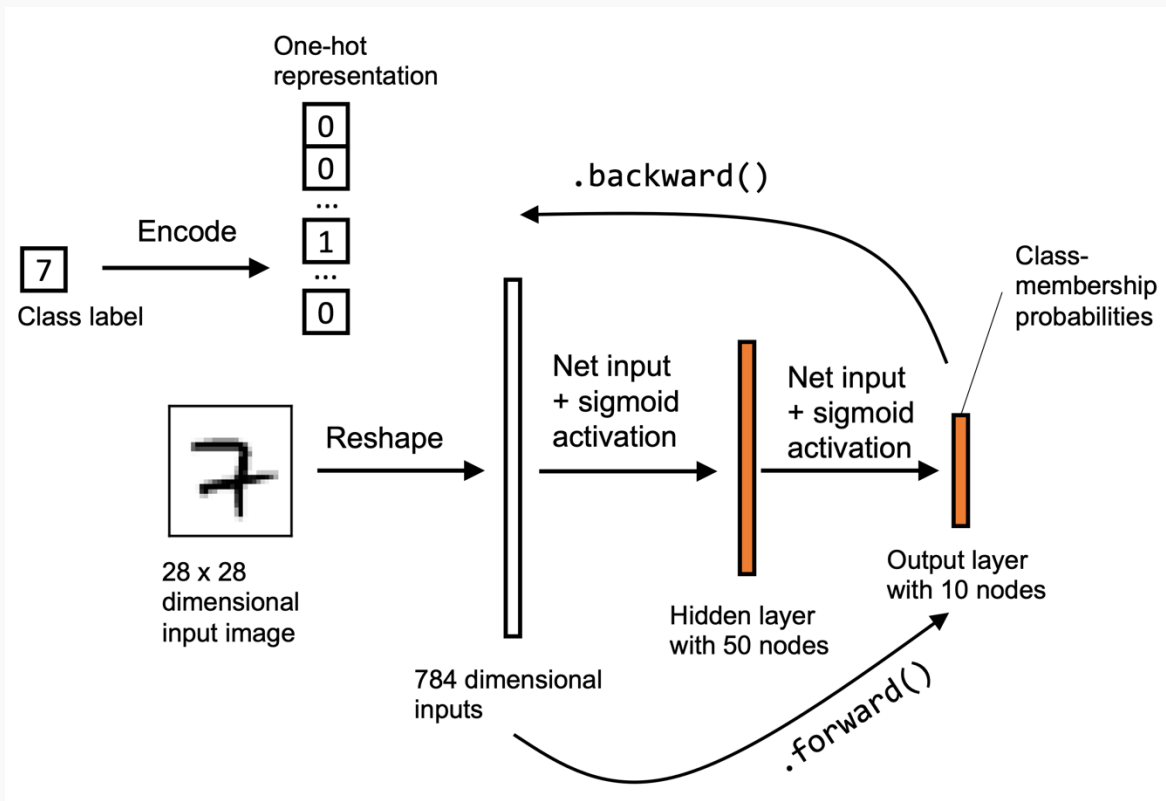
*Hand-written digit recognition: MNIST data*



# MNIST NN

*Today we'll work on forward algorithm..*

*Next lecture backpropagation, gradient descent, training, and evaluating network performance*



## MNIST NN utility code

*How would the following sample labels be one-hot encoded: [1, 0, 2] ?*

```
def sigmoid(z):  
    return 1. / (1. + np.exp(-z))  
  
def int_to_onehot(y, num_labels):  
  
    ary = np.zeros((y.shape[0], num_labels))  
    for i, val in enumerate(y):  
        ary[i, val] = 1  
  
    return ary
```

60

# MNIST NN initialization

## *Initialize weights and biases*

```
class NeuralNetMLP:

    def __init__(self, num_features, num_hidden, num_classes, random_seed=123):
        super().__init__()

        self.num_classes = num_classes

        # hidden
        rng = np.random.RandomState(random_seed)

        self.weight_h = rng.normal(
            loc=0.0, scale=0.1, size=(num_hidden, num_features))
        self.bias_h = np.zeros(num_hidden)

        # output
        self.weight_out = rng.normal(
            loc=0.0, scale=0.1, size=(num_classes, num_hidden))
        self.bias_out = np.zeros(num_classes)
```

# Forward Algorithm

---

**Require:** Network depth,  $l$

**Require:**  $\mathbf{W}^{(i)}, i \in \{1, \dots, l\}$ , the weight matrices of the model

**Require:**  $\mathbf{b}^{(i)}, i \in \{1, \dots, l\}$ , the bias parameters of the model

**Require:**  $\mathbf{x}$ , the input to process

**Require:**  $\mathbf{y}$ , the target output

$$\mathbf{h}^{(0)} = \mathbf{x}$$

**for**  $k = 1, \dots, l$  **do**

$$\mathbf{a}^{(k)} = \mathbf{b}^{(k)} + \mathbf{W}^{(k)}\mathbf{h}^{(k-1)}$$

$$\mathbf{h}^{(k)} = f(\mathbf{a}^{(k)})$$

**end for**

$$\hat{\mathbf{y}} = \mathbf{h}^{(l)}$$

$$J = L(\hat{\mathbf{y}}, \mathbf{y}) + \lambda\Omega(\theta)$$

---

# MNIST NN Forward

*How do we use the output for prediction?*

```
def forward(self, x):  
    # Hidden layer  
    # input dim: [n_examples, n_features] dot [n_hidden, n_features].T  
    # output dim: [n_examples, n_hidden]  
    z_h = np.dot(x, self.weight_h.T) + self.bias_h  
    a_h = sigmoid(z_h)  
  
    # Output layer  
    # input dim: [n_examples, n_hidden] dot [n_classes, n_hidden].T  
    # output dim: [n_examples, n_classes]  
    z_out = np.dot(a_h, self.weight_out.T) + self.bias_out  
    a_out = sigmoid(z_out)  
    return a_h, a_out
```

# Today's Learning Objectives

Students will be able to:

- ✓ Review: Describe what a neural network (NN) is
- ✓ Apply the forward algorithm on the MNIST NN
  - Trace Backpropagation on a small NN
  - Understand the role of gradient descent in training NN

# Loss Functions

---

*As with other models, we have choice of loss functions, such as*

- 1. MSE*
- 2. Negative Log Likelihood (Binary Cross Entropy) or its generalization for multi-class classification*

*We'll look at MSE for MNIT*



## How to calculate MSE for MNIST

---

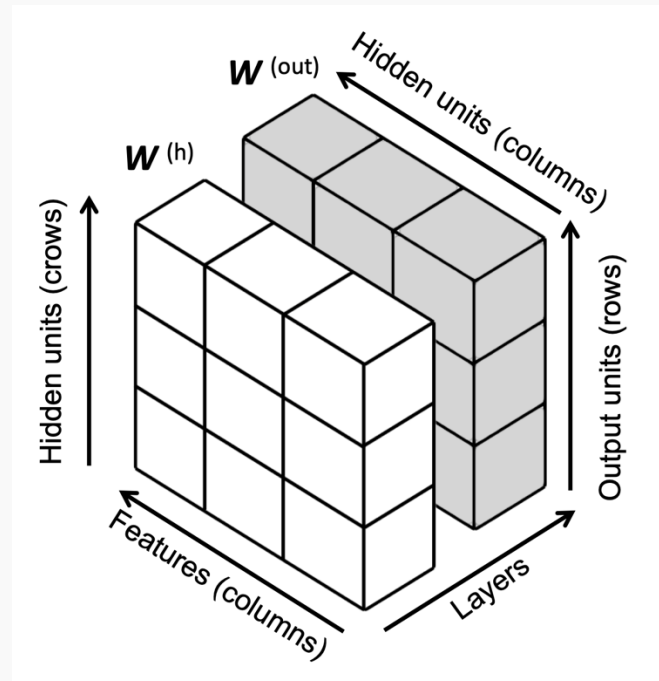
*We have a vector of  $y$  labels that are one-hot encoded.*

*If image is 2:  $[0, 0, 1, 0, 0, 0, 0, 0, 0]$*

*Output from NN =  $[.1, .03, .9, .1, .1, .1, .1, .3, .1, .1]$*

## Minimize Loss

*As with any model, we want to minimize our loss function.  
To do so we need to calculate partial derivative of our loss function with respect to all the weights in the model.*



## Applying the Chain Rule

---

$$\frac{d}{dx}[f(g(x))] = \frac{df}{dg} \cdot \frac{dg}{dx}$$

$$\frac{d}{dx}[f(g(h(u(v(x)))))] = \frac{df}{dg} \cdot \frac{dg}{dh} \cdot \frac{dh}{du} \cdot \frac{du}{dv} \cdot \frac{dv}{dx}$$

# Backprop Algorithm

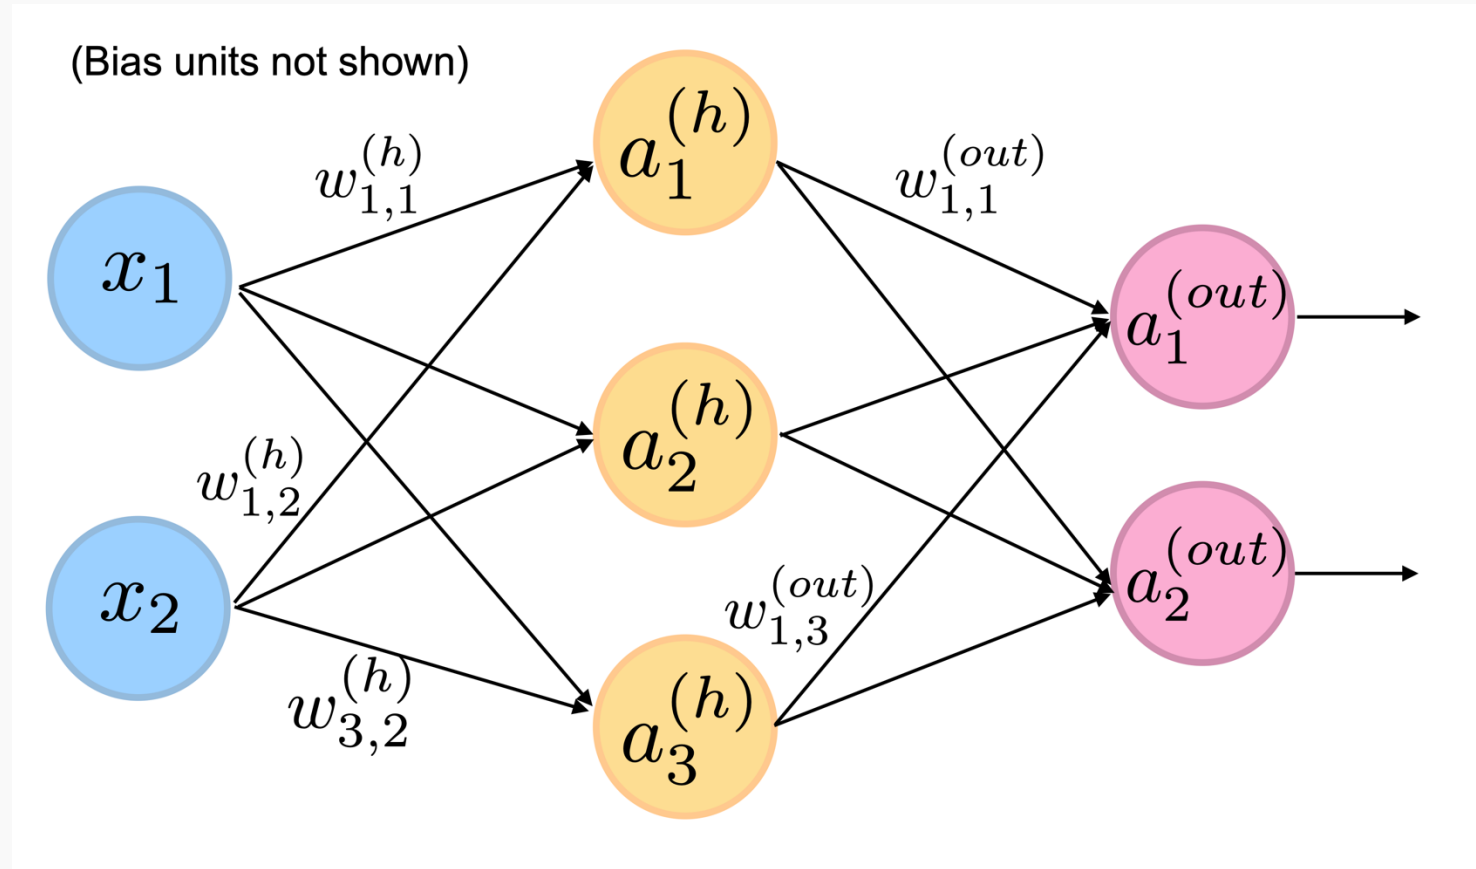
---

1. *Compute the gradient of loss function with respect to output layer*
2. *For each layer*
  - i. *Convert gradient on layer's output into a gradient with respect to the net input before activation*
  - ii. *Compute gradients on weights and biases by using gradient of the net input with respect to weights*
  - iii. *Propagate the gradients backwards in NN (i.e. save to reuse for gradients in lower layers)*

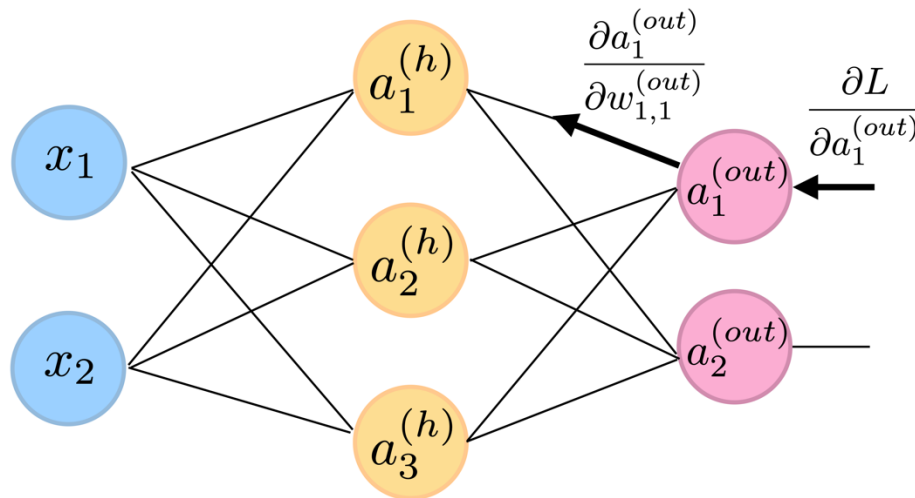
***Gradients give us how much output layer or weight should change  
To reduce loss.***

60

# Tracing Backprop



# Tracing Backprop



Gradient for output layer weight:

$$\frac{\partial L}{\partial w_{1,1}^{(out)}} = \frac{\partial L}{\partial a_1^{(out)}} \cdot \frac{\partial a_1^{(out)}}{\partial w_{1,1}^{(out)}}$$

$$\frac{\partial L}{\partial w_{1,1}^{(out)}} = \frac{\partial L}{\partial a_1^{(out)}} \cdot \frac{\partial a_1^{(out)}}{\partial z_1^{(out)}} \cdot \frac{\partial z_1^{(out)}}{\partial w_{1,1}^{(out)}}$$

## Calculating the partial derivatives

---

$$\frac{\partial L}{\partial w_{1,1}^{(out)}} = \frac{\partial L}{\partial a_1^{(out)}} \cdot \frac{\partial a_1^{(out)}}{\partial z_1^{(out)}} \cdot \frac{\partial z_1^{(out)}}{\partial w_{1,1}^{(out)}}$$

## Calculating the partial derivatives

---

$$\frac{\partial L}{\partial w_{1,1}^{(out)}} = \frac{\partial L}{\partial a_1^{(out)}} \cdot \frac{\partial a_1^{(out)}}{\partial z_1^{(out)}} \cdot \frac{\partial z_1^{(out)}}{\partial w_{1,1}^{(out)}}$$



## Calculating the partial derivatives

---

$$\frac{\partial L}{\partial w_{1,1}^{(out)}} = \frac{\partial L}{\partial a_1^{(out)}} \cdot \frac{\partial a_1^{(out)}}{\partial z_1^{(out)}} \cdot \frac{\partial z_1^{(out)}}{\partial w_{1,1}^{(out)}}$$

## Calculating the partial derivatives

---

$$\frac{\partial L}{\partial w_{1,1}^{(out)}} = \frac{\partial L}{\partial a_1^{(out)}} \cdot \frac{\partial a_1^{(out)}}{\partial z_1^{(out)}} \cdot \frac{\partial z_1^{(out)}}{\partial w_{1,1}^{(out)}}$$

## Calculating the partial derivatives

---

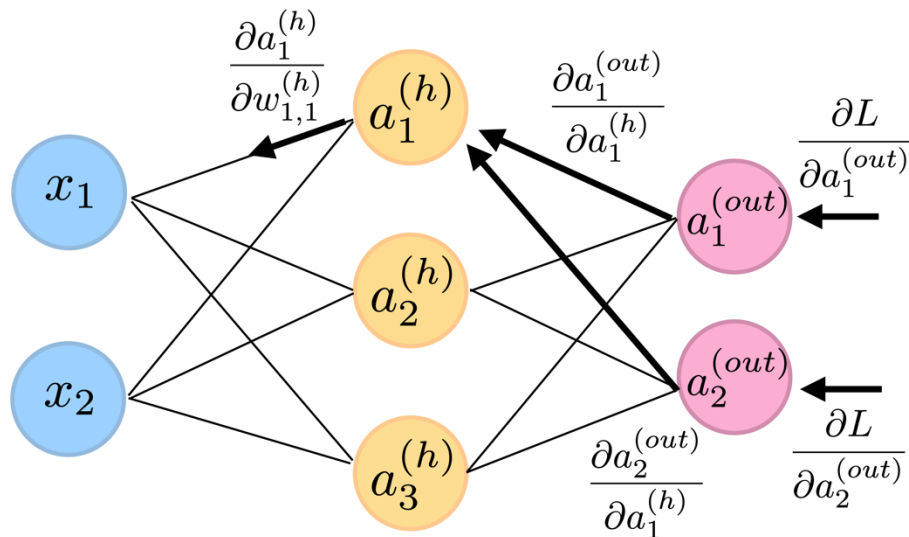
$$\frac{\partial L}{\partial w_{1,1}^{(out)}} = \frac{\partial L}{\partial a_1^{(out)}} \cdot \frac{\partial a_1^{(out)}}{\partial z_1^{(out)}} \cdot \frac{\partial z_1^{(out)}}{\partial w_{1,1}^{(out)}}$$

***Update weights for gradient descent***

***Reuse gradient of loss with respect to net output***

60

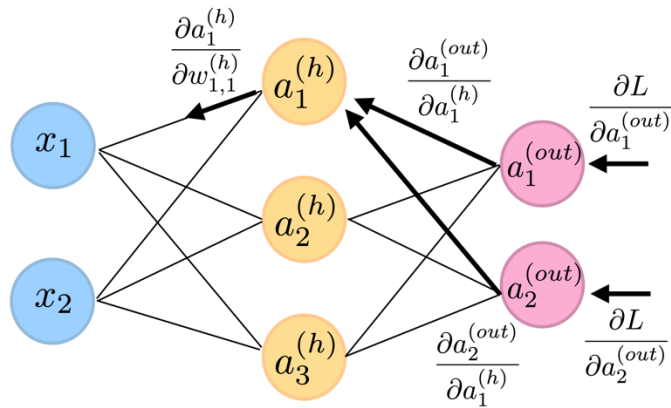
# Tracing Backprop



Gradient for hidden layer weight:

$$\begin{aligned} \frac{\partial L}{\partial w_{1,1}^{(h)}} &= \frac{\partial L}{\partial a_1^{(out)}} \cdot \frac{\partial a_1^{(out)}}{\partial a_1^{(h)}} \cdot \frac{\partial a_1^{(h)}}{\partial w_{1,1}^{(h)}} \\ &+ \frac{\partial L}{\partial a_2^{(out)}} \cdot \frac{\partial a_2^{(out)}}{\partial a_1^{(h)}} \cdot \frac{\partial a_1^{(h)}}{\partial w_{1,1}^{(h)}} \end{aligned}$$

# Tracing Backprop



Gradient for hidden layer weight:

$$\begin{aligned} \frac{\partial L}{\partial w_{1,1}^{(h)}} &= \frac{\partial L}{\partial a_1^{(out)}} \cdot \frac{\partial a_1^{(out)}}{\partial a_1^{(h)}} \cdot \frac{\partial a_1^{(h)}}{\partial w_{1,1}^{(h)}} \\ &+ \frac{\partial L}{\partial a_2^{(out)}} \cdot \frac{\partial a_2^{(out)}}{\partial a_1^{(h)}} \cdot \frac{\partial a_1^{(h)}}{\partial w_{1,1}^{(h)}} \end{aligned}$$

# Today's Learning Objectives

Students will be able to:

- ✓ Review: Describe what a neural network (NN) is
- ✓ Apply the forward algorithm on the MNIST NN
- ✓ Trace Backpropagation on a small NN
  - Understand the role of gradient descent in training NN

# Training to minimize the loss function

---

**Question:** What is the mathematical function that describes the slope?

**Derivative**

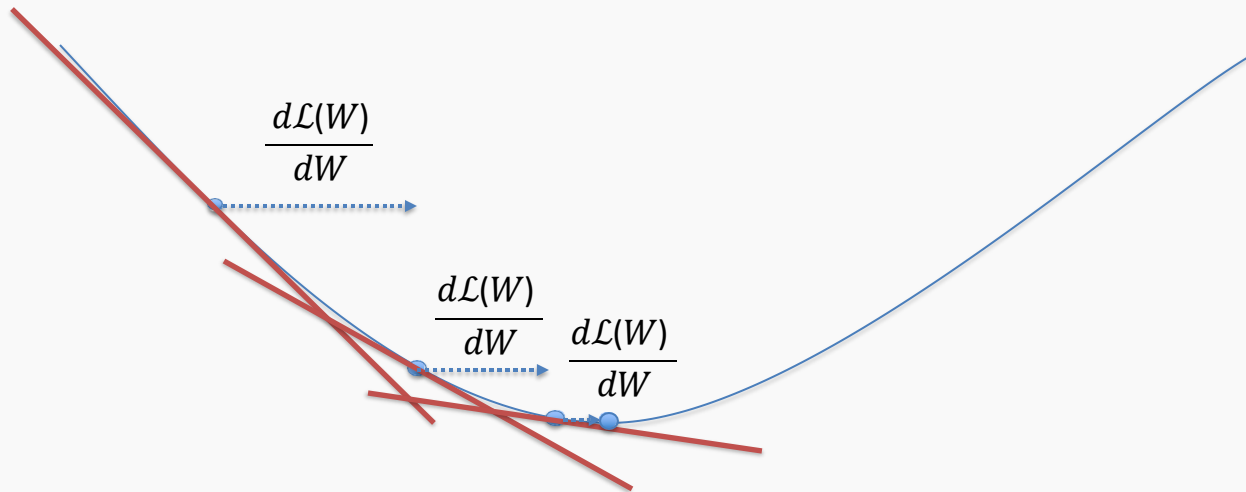
**Question:** How do we generalize this to more than one predictor?

**Take the derivative with respect to each coefficient and do the same sequentially**

**Question:** What do you think is a good approach for telling the model how to change (what is the step size) to become better?

# Gradient Descent

If the step is proportional to the slope then you avoid overshooting the minimum.  
How?





# Backprop Algorithm

---

1. *Compute the gradient of loss function with respect to output layer*
2. *For each layer*
  - i. *Convert gradient on layer's output into a gradient with respect to the net input before activation*
  - ii. *Compute gradients on weights and biases by using gradient of the net input with respect to weights*
  - iii. *Propagate the gradients backwards in NN (i.e. save to reuse for gradients in lower layers)*

***Gradients give us how much output layer or weight should change  
To reduce loss.***

60

# Gradient Descent

---

1. *Initialize small weights*
2. *For each sample*
  - i. *Apply the forward algorithm*
  - ii. *Apply backprop*
  - iii. *Use the gradients of the weights and biases to update the weights and biases*

*Repeat step 2 for a number of **epochs**, complete pass through Training data*

## Updating the weights

---

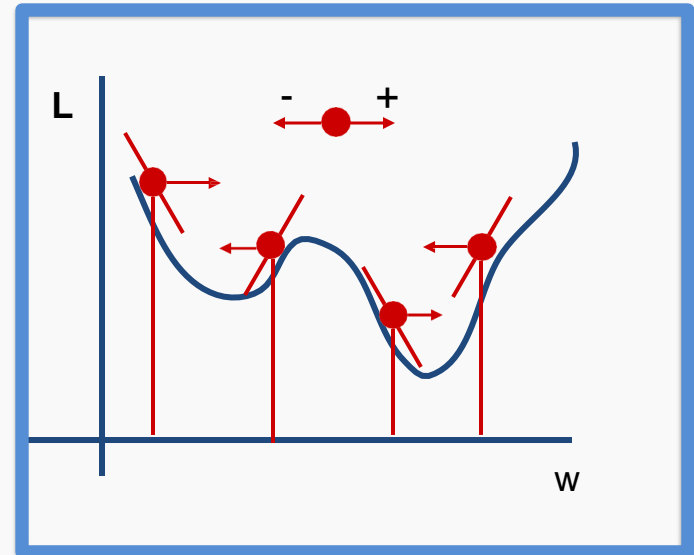
*For each weight*

$$\Delta w = \frac{\partial L}{\partial w}$$

$$w^{\text{new}} = w^{\text{old}} - \eta \Delta w$$

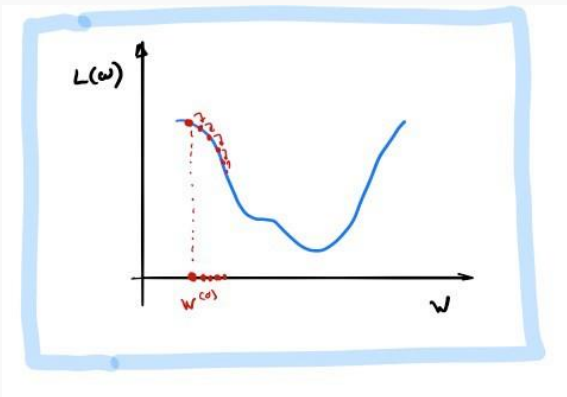
# Gradient Descent

- Algorithm for optimization of first order to finding a minimum of a function.
- It is an iterative method.
- $L$  is decreasing much faster in the direction of the negative derivative.
- The learning rate is controlled by the magnitude of  $\eta$ .

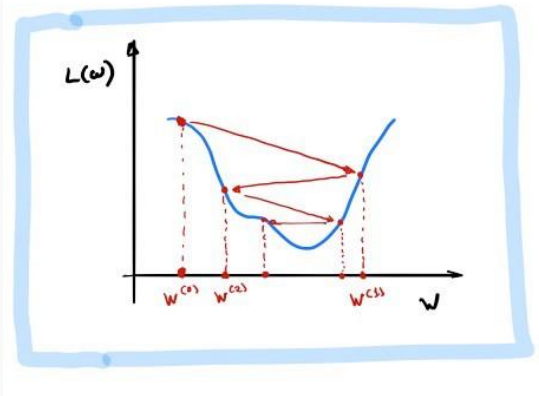


# Learning Rate

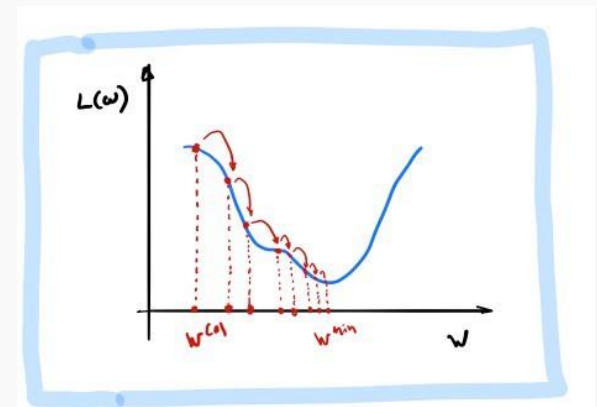
Our choice of the learning rate has a significant impact on the performance of gradient descent.



When  $\eta$  is too small, the algorithm makes very little progress.



When  $\eta$  is too large, the algorithm may overshoot the minimum and has crazy oscillations.



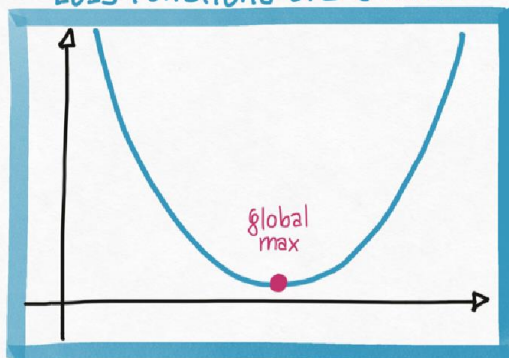
When  $\eta$  is appropriate, the algorithm will find the minimum. The algorithm **converges**!

# Local vs Global Minima

If we choose  $\eta$  correctly, then gradient descent will converge to a stationary point. But will this point be a **global minimum**?

If the function is convex then the stationary point will be a global minimum.

Linear & Polynomial Regression  
Loss Functions are Convex



Hessian (2nd Derivative) positive semi-definite everywhere.

Every stationary point of the gradient is a global min.

Neural Network Regression Loss  
Functions are not Convex



Neural networks with different weights can correspond to the same function.

Most stationary points are local minima but not global optima.

---

## Batch and Stochastic Gradient Descent

Instead of using all the training examples for every step,

For **batch gradient descent**,

In each epoch use a subset or **batch** of samples. (These can be randomly selected). Use average of the gradients from the batches to update the weights

For **stochastic gradient descent**,

Use a randomly selected sample to update the weights a sample at a time. (Like batch, but here size = 1)

DATA



↓  
calculate  $L_1 \Rightarrow \frac{\partial L_1}{\partial w} \Rightarrow w^* = w - \eta \frac{\partial L_1}{\partial w}$

↓  
calculate  $L_2 \Rightarrow \frac{\partial L_2}{\partial w} \Rightarrow w^* = w - \eta \frac{\partial L_2}{\partial w}$

⋮  
↓  
 $L_B \Rightarrow \frac{\partial L_B}{\partial w} \Rightarrow w^* = w - \eta \frac{\partial L_B}{\partial w}$

COMPLETE DATA  $\Rightarrow$  ONE EPOCH

RESHUFFLE DATA AND REPEAT



# Your turn: MNIST NN

Please get the Jupyter notebook for GPU data:

Go to Notebook (also on BruinLearn):

[https://colab.research.google.com/drive/10BkOu\\_bVBs3af2NXFrD48iK8s3MJIScj?usp=sharing](https://colab.research.google.com/drive/10BkOu_bVBs3af2NXFrD48iK8s3MJIScj?usp=sharing)

Save a copy to your Google Drive and keep notes there...

# Today's Learning Objectives

Students will be able to:

- ✓ Review: Describe what a neural network (NN) is
- ✓ Apply the forward algorithm on the MNIST NN
- ✓ Trace Backpropagation on a small NN
- ✓ Understand the role of gradient descent in training NN

*Citations:.*

Goodfellow, I., Bengio, Y., Courville, A. (2016). *Deep Learning*. MIT Press.

Sebastian **Raschka**, Yuxi (Hayden) **Liu**, and Vahid Mirjalili. **Machine Learning** with PyTorch and Scikit-Learn. Packt Publishing, **2022**.

*Baharan Mirzasoleiman, UCLA CS M148 Winter 2024 Lecture 13 Notes*

# Thank You

---